

CSE 421/521 - Operating Systems Fall 2018

LECTURE - XVIII

FILE SYSTEMS - I

Tevfik Koşar

University at Buffalo
November 1st, 2018

Roadmap

- File Allocation Methods
 - Contiguous
 - Linked
 - Indexed
- Free Space Management
 - Bit vector
 - Linked List
 - Grouping
 - Counting



File Allocation Methods


- An allocation method refers to **how disk blocks are allocated & organized for files on the disk:**
- Contiguous allocation
- Linked allocation
- Indexed allocation

1. Contiguous Allocation

An allocation method refers to how disk blocks are allocated for files:

- ✦ **Contiguous allocation** - each file occupies set of contiguous blocks
- **Best performance** in most cases (both sequential and random access)
- **Simple** - only starting location (block #) and length (number of blocks) are required
- **Problems** include finding space for file, knowing file size,


Contiguous Allocation of Disk Space



Contiguous Allocation Address Translation

✦ Mapping from logical to physical disk address, given:

- LA: logical address
- 512: disk block size
- Q: quotient
- R: remainder




Block to be accessed = Q + starting address

Displacement into block = R


2. Linked Allocation

- Each file is a linked list of disk blocks: blocks may be scattered anywhere on the disk.




- + Simple - need only starting address
- + Free-space management system - no waste of space
- + Defragmentation not necessary for file allocation
- No random access, locating a block can take many disk seek & I/O
- Extra space required for pointers
- Reliability: what if a pointer gets corrupted?

Linked Allocation



Linked Allocation

- ✦ Mapping (Assuming pointer takes only **one byte**)




Block to be accessed is the Qth block in the linked chain of blocks representing the file.

Displacement into block = $R + 1$


File-Allocation Table

- ✦ FAT (File Allocation Table) variation
 - Beginning of volume has table, indexed by block number
 - Much like a linked list, but faster on disk and cacheable
 - New block allocation simple



3. Indexed Allocation

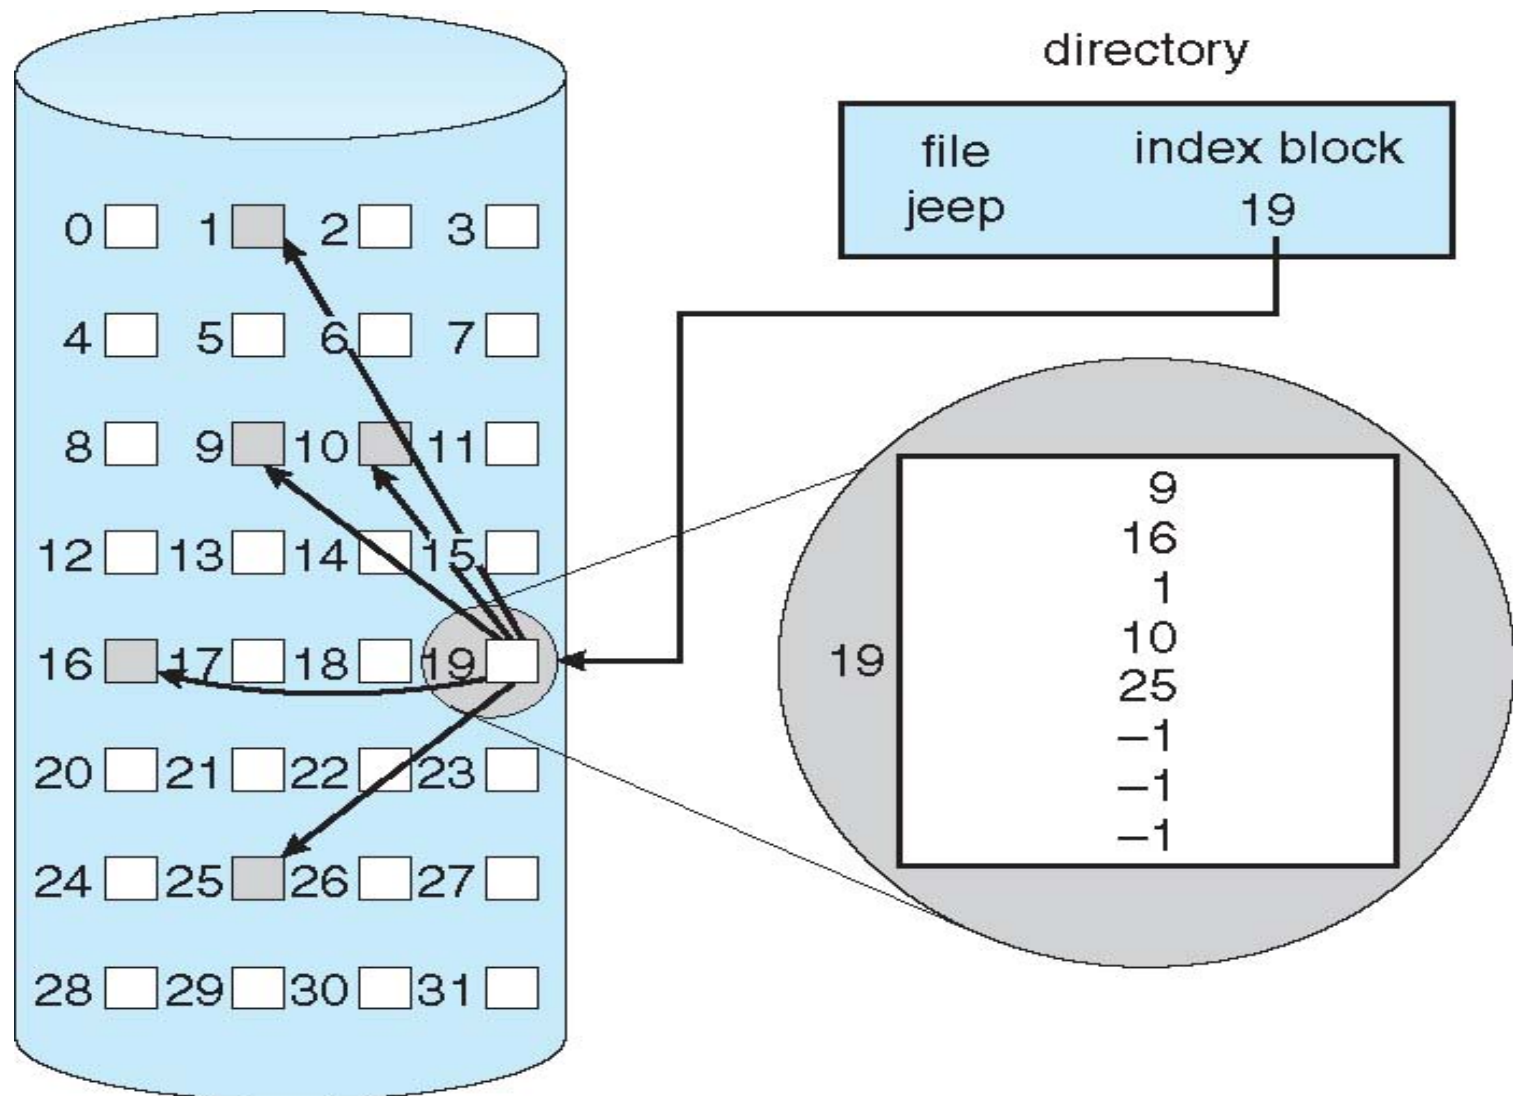
- Brings all pointers together into the *index block*, to allow random access to file blocks.
- Logical view.



index table


- + Supports direct access
- + Prevents external fragmentation
- High pointer overhead --> wasted space

Example of Indexed Allocation



Indexed Allocation

- ◆ Need index table
- ◆ Random access
- ◆ Dynamic access without external fragmentation, but have overhead of index block
- ◆ Mapping from logical to physical in a file of maximum size of 256K bytes and block size of 512 bytes. We need only 1 block for index table



Q = displacement into index table

R = displacement into block

Performance

- ✦ Best method depends on file access type
 - **Contiguous** great for sequential and random
- ✦ **Linked** good for sequential, not random
- ✦ Declare access type at creation -> select either contiguous or linked
- ✦ **Indexed** more complex
 - Single block access could require 2 index block reads then data block read
 - Clustering can help improve throughput, reduce CPU overhead

Exercise

Consider a file system on a disk that has both logical and physical **block sizes of 512 bytes**. Assume that the information about each file is already in memory. For each of the three allocation strategies (contiguous, linked, and indexed), answer these questions:

- a. How is the **logical-to-physical address mapping** accomplished in this system? (For the indexed allocation, assume that a file is always less than 512 blocks long.)
- b. If we are **currently at the 10th logical block** of the file and **want to access the 4th logical block**. How many physical blocks must be read from the disk?

PS: Assume a pointer needs only one byte space.

Exercise - Solution

a. Contiguous:

Divide the logical address by 512 with X and Y the resulting quotient and remainder respectively.

1. Add X to Z to obtain the physical block number. Y is the displacement into that block.
2. One.

Exercise - Solution

b. Linked:

Divide the logical physical address by 511 with X and Y the resulting quotient and remainder respectively.

1. Chase down the linked list (getting $X + 1$ blocks). $Y + 1$ is the displacement into the last physical block.
2. Four.

Exercise - Solution

c. Indexed:

Divide the logical address by 512 with X and Y the resulting quotient and remainder respectively.

1. Get the index block into memory. Physical block address is contained in the index block at location X . Y is the displacement into the desired physical block.


2. Two

Free Space Management

- Disk space limited
- Need to re-use the space from deleted files
- To keep track of free disk space, the system maintains a **free-space list**
 - Records all free disk blocks
- Implemented using
 - Bit vectors
 - Linked lists


Free-Space Management (Cont.)

- Linked List Approach



Free-Space Management (Cont.)

- Bit vector (n blocks)
 - Each block is represented by 1 bit
 - 1: free, 0: allocated




$$\text{bit}[i] = \begin{cases} 1 \Rightarrow \text{block}[i] \text{ free} \\ 0 \Rightarrow \text{block}[i] \text{ occupied} \end{cases}$$

- e.g. 0000111110001000100010000

Free-Space Management (Cont.)

- Bit Vector Approach




Free-Space Management (Cont.)

- Bit map requires extra space
 - Example:
 block size = 2^{12} bytes
 disk size = 2^{30} bytes (1 gigabyte)
 $n = 2^{30}/2^{12} = 2^{18}$ bits (or 32K bytes)
- Easy to get contiguous files
- **Linked list (free list)**
 - Cannot get contiguous space easily
 - requires substantial I/O
- **Grouping**
 - Modification of free-list
 - Store addresses of n free blocks in the first free block
- **Counting**
 - Rather than keeping list of n free addresses:
 - Keep the address of the first free block
 - And the number n of free contiguous blocks that follow it

Free-Space Management (Cont.)

Bit Map/Linked List/Grouping/Counting

free-list head



grouping ($n=3$)

1 2,3,7

7 8,9,12

12 14,15,16

16 17,18,-1


bit map: 011100011100101111100000

counting: (1,3), (7, 3), (12, 1), (14, 5)

os12

27

Exercise - 1



Given the current file allocation on the disk in this figure, show the linked list, bit-map, counting, and grouping (n=4) representations of the free block list.

Exercise - 2

- In terms of reliability and performance, compare bit vector implementation of a free block list with keeping a list of free blocks where the first few bytes of each free block provide the logical sector number of the next free block.


Exercise - 2 Solution

performance: Bit vector implementation is more efficient since it allows fast and random access to free blocks, linked list approach allows only sequential access, and each access results in a disk read. Bit vector implementation can also find n consecutive free blocks much faster. (Assuming bit vector is kept in memory)

reliability: If an item in a linked list is lost, you cannot access the rest of the list. With a bit vector, only those items are lost. Also, it's possible to have multiple copies of the bit vector since it is a more compact representation. Although keeping the bit vector in memory seems to be unreliable, you can always keep an extra copy on the disk.

Any Questions?

- File Allocation Methods
 - Contiguous
 - Linked
 - Indexed
- Free Space Management
 - Bit vector
 - Linked List
 - Grouping
 - Counting



Reading assignment: Chapter 12 from Silberschatz

Acknowledgements

- “Operating Systems Concepts” book and supplementary material by A. Silberschatz, P. Galvin and G. Gagne
- “Operating Systems: Internals and Design Principles” book and supplementary material by W. Stallings
- “Modern Operating Systems” book and supplementary material by A. Tanenbaum
- R. Doursat and M. Yuksel from UNR